

Stat243: Problem Set 5

Linqing Wei

October 17, 2017

Problem 2

$$\begin{aligned}1 &= 1.0 * 2^{1023-1023} \\2 &= 1.0 * 2^{1024-1023} \\3 &= 1.5 * 2^{1024-1023} \\4 &= 1.0 * 2^{1025-1023} \\2^{53} - 2 &= 1.0 * (2^{1076-1023}) - 1.0 * (2^{1024-1023}) \\2^{53} &= 1.0 * (2^{1076-1023}) \\2^{53} + 1 &= 1.0 * (2^{1076-1023}) + 1.0 * (2^{1023-1023}) \\2^{53} + 2 &= 1.0 * (2^{1076-1023}) + 1.0 * (2^{1024-1023})\end{aligned}$$

```
> library(pryr)
> 2^53

[1] 9.007199e+15

> 2^53 + 1

[1] 9.007199e+15

> 2^53 + 2

[1] 9.007199e+15

> bits(2^53) == bits(2^53 + 1 )

[1] TRUE

> bits(2^53) == bits(2^53 + 2)

[1] FALSE
```

Since the bits of 2^{53} are equal to the bits of $2^{53} + 1$, it means $2^{53} + 1$ cannot be represented exactly. Since the bits of 2^{53} are not equal to the bits of $2^{53} + 2$, it means $2^{53} + 2$ is unique, CAN be represented exactly.

```
> 2^53
[1] 9.007199e+15
> 2^53 - 1
[1] 9.007199e+15
> 2^53 + 1
[1] 9.007199e+15
> bits(2^53) == bits(2^53 - 1)
[1] FALSE
> bits (2^53) == bits(2^53 + 1)
[1] TRUE
```

2^{53} can be represented exactly, and $2^{53} - 1$ can be represented exactly since it's less than 2^{53} . The bits of 2^{53} are equal to $2^{53} + 1$, meaning that the bits of $2^{53} + 1$ are not unique, CANNOT be represented exactly. Therefore, the spacing is 2 of numbers starting with 2^{53}

```
> 2^54
[1] 1.80144e+16
> 2^54 - 2
[1] 1.80144e+16
> 2^54 + 2
[1] 1.80144e+16
> bits(2^54) == bits (2^54-2)
[1] FALSE
> bits(2^54) == bits (2^54+2)
[1] TRUE
```

When analyzing the spacing of 2^{54} , let's assume 2^{54} is fixed and can be represented exactly. Since the bits of 2^{54} is not equal to $2^{54} - 2$, meaning that the bits of $2^{54} - 2$ is unique, CAN be represented exactly. However, the bits of 2^{54} is equal to $2^{54} + 2$, meaning that the bits of $2^{54} + 2$ is not unique, CANNOT be represented exactly. Therefore, based on the representation of $2^{54} - 2$ and $2^{54} + 2$, spacing of numbers starting with 2^{54} is 4.

Problem 3

3a

```
> library(microbenchmark)
> library(data.table)
```

This data.table install has not detected OpenMP support. It will work but slower in single t

```
> a = as.integer(rnorm(1e8))
> b = as.numeric(rnorm(1e8))
> system.time(c <- copy(a)) #Make a deep copy of a
```

```
      user  system elapsed
0.114    0.109    0.223
```

```
> system.time(d <- copy(b)) #Make a deep copy of b
```

```
      user  system elapsed
0.268    0.239    0.507
```

It is faster to copy a large vector of integers than to copy a numeric vector, since each integer takes 4 bytes while each numeric value(double) takes 8 bytes.

3b

```
> I = as.integer(rnorm(1e6))
> N = as.numeric(rnorm(1e6))
> microbenchmark(a[c(1:length(I)/2)])
```

Unit: milliseconds

	expr	min	lq	mean	median	uq	max
a[c(1:length(I)/2)]		22.45112	32.75253	38.06377	34.07807	36.23189	170.5435
neval							
100							

```
> microbenchmark(b[c(1:length(N)/2)])
```

Unit: milliseconds

	expr	min	lq	mean	median	uq	max
b[c(1:length(N)/2)]		22.88178	34.41606	39.24622	35.61536	37.20703	173.7124
neval							
100							

It is a bit faster to take a subset of size $k = n/2$ from an integer vector than from a numeric vector, since each interger takes 4 bytes while each numeric value(double) takes 8 bytes. However, the running time does not show much a difference.

Problem 4

4a

The naive algorithm of running matrix column by column takes at least $O(n^2)$ times of running. The parallel computing algorithm takes $O(np)$ times for partial matrix multiplication. Then summing up the partial results would take roughly $O(\log n)$ running times, which is more efficient.

4b

Approach A is better for minimizing communication. Each time submatrix of Y interacts with the whole matrix X. The total communication times are p (or j , if you think about it in terms of number of tasks). However, each time, the whole matrix of x is passed into the function such that the memory usage is large for approach A, with at least $p \cdot \text{objectsize}(x)$ amount of memory.

Approach B is better for minimizing memory usage. Each time, only a subset of X and a subset of Y are passed into the function. However, the communication cost is high since the total number of partitioned blocks is large. It would take at least p^2 of communication times.

Extra Credit

R uses IEEE754 scientific notation. For example, 0.1 has an index of -4, 0.2 has an index of -3 and 0.3 has an index of -2.

When summing up 0.1 and 0.2, R will change the index of 0.1 from -4 to -3 (raising the power), which leads to a change in floating points.

Originally:

$$0.1 = 1.1001100110011001100110011001100110011001100110011010.$$

After change,

$$0.1 = 0.1100110011001100110011001100110011001100110011001101.$$

After adjusting the floating point,

$$0.1 + 0.2 = 1.00110011001100110011001100110011001100110011001100111,$$

which is different from the original representation of 0.3.

However, $0.2 + 0.3 == 0.5$ is true because when doing addition using IEEE754 scientific notation, we unify the index (power) of 0.2 and 0.3 as -2.

After performing the similar addition as in "0.1 + 0.2" case, and readjust the floating point position, the final result of 0.5 is actually equal to the real representation of 0.5.